

# Rohan Bhide

rohanbhide19@gmail.com | linkedin.com/in/rohanbhide | github.com/rhn19

## EDUCATION

### University of California-San Diego (UCSD)

Master of Science (MS) - Computer Science; GPA 3.9/4.0

San Diego, CA

Sep 2022 – Dec 2023

### Pune Institute of Computer Technology (PICT)

Bachelor of Engineering (BE) - Computer Engineering; GPA 9.67/10.0

Pune, India

Jun 2018 – May 2022

## EXPERIENCE

### Sony Electronics

San Diego, CA

Sr. Software Engineer

Mar 2024 – Present

- Integrated sliding-window KV cache support into ExecuTorch's Vulkan backend via a fused indexed cache update GLSL shader, profiling against XNNPack for latency, memory, thermals on Adreno SM8850 and Mali-G715.
- Achieved 92% function-calling accuracy on 4bit/8bit quantized Gemma/Llama by fine-tuning with QAT and LoRA.
- Expanded 3D ecosystem for Spatial Reality Displays (SR1/SR2) by developing Sony's first C++ OpenXR runtime, enabling support for 20+ professional applications including Omniverse, VRED, Blender, and TwinMotion.
- Ensured sub-16ms motion-to-photon latency by engineering a multi-threaded mirror window for DirectX/OpenGL interop via shared handles and synchronization fences, integrating Nvidia Reflex for performance measurement.
- Enabled high-performance LLM inference on devices with less than 2GB RAM by pruning and distilling Llama3 to sub-1B parameters, facilitating CPU-only execution integrated with Whisper and LangChain for embedded RAG.
- Secured a RAG-based troubleshooting assistant against adversarial threats, maintaining <10% successful attacks on high severity tests by implementing Promptfoo red-teaming and token-optimized prompt engineering.
- Scaled creative insights to 150K+ professionals by architecting a serverless AWS platform and Android app leveraging multimodal Gemini models for real-time composition analysis and on-device framing guidance.

VR Software Engineering Intern

Oct 2023 – Mar 2024

- Enabled native, glasses-free 3D previewing within the Maya viewport by developing a C++ plugin that captures dual real-time camera perspectives using off-axis projection for Spatial Reality Displays (SR1/SR2).
- Improved spatial accuracy and framing efficiency for 3D artists by implementing viewport overrides and custom hardware-representative gizmos that constrain camera frustums to the gizmo's physical bounds.
- Ensured seamless UI/UX consistency across 8+ varied unit and scale configurations by creating a MEL-based management interface that automates gizmo lifecycle and alignment logic.
- Reduced deployment overhead by 60% by architecting an Azure DevOps CI/CD pipeline utilizing Visual Studio and WiX Toolset to automate compilation, signed installer generation, and delivery of DLLs, shaders, and scripts.

### The Design Lab @UCSD

San Diego, CA

Graduate Researcher

Oct 2022 – Dec 2023

- Advanced driving safety research by programmatically generating 26 complex route configurations in Python for DReyeVR (CARLA/Unreal Engine), modeled after federal crash data to simulate high-risk scenarios.
- Engineered high-fidelity multiplayer Unity FPS with NavMesh AI pathfinding and raycast combat, incorporating diegetic UI for emotion visualization to achieve a 150% increase in self-reported social presence scores.
- Quantified impact of affective projection on virtual social presence through a user study (N=18) by integrating AffectNet facial expression recognition via low-latency UDP sockets. [escholarship.org/uc/item/2xv7w1zq]

### AutoVRse

Bengaluru, India

AR/VR Software Engineering Intern

Feb 2022 – Jul 2022

- Built an industrial assembly simulation using VRTK by leading 3 developers, overseeing the Git workflow, and configuring Unity SmartMerge to resolve complex scene conflicts.
- Reduced enterprise training time by 75% by developing a VR power tool simulation for Meta Quest featuring C# voice-guided lessons and HLSL shaders for pulsing spatial highlights.
- Engineered a cooperative multiplayer skating system via 6-DOF controller velocity translation and hosted 200+ users in VRChat by deploying network-optimized object synchronization for shared interactable items using Udon.

### Persistent Systems

Pune, India

AR/VR Software Engineering Intern

Jun 2021 – Aug 2021

- Streamlined consumer networking hardware installation via a Unity/ARFoundation app, using image tracking to overlay dynamic 3D guides for 7 components (incl. LAN, Power, RJ-11) synced with troubleshooting instructions.
- Enabled dynamic instruction delivery and remote feedback management via Firebase Realtime Database REST APIs; improved accessibility by developing text-to-speech, voice commands, and interactive modules.

## Computational Linguistics Lab @PICT

Pune, India

*Undergraduate Researcher*

*Sep 2020 – Aug 2021*

- Architected end-to-end Tesseract and OpenCV OCR pipeline to digitize 20+ textbooks into structured formats.
- Created a dataset of 4000+ textbook questions and achieved 88% Bloom's Taxonomy classification accuracy by fine-tuning BERT with a multi-label classification head in PyTorch.
- Reduced manual sourcing time by 50% using Matplotlib/Seaborn heatmaps for cognitive complexity distribution.

## iMocha

Pune, India

*Natural Language Processing Intern*

*Mar 2020 – Sep 2020*

- Achieved 81% accuracy evaluating candidate essay grammar, tone, and relevance by fine-tuning BERT in PyTorch and integrating the inference pipeline into a Django backend with robust error-handling.
- Automated analysis of candidate fluency and pronunciation for 1,000+ interviews via Azure's Speech-to-Text APIs.

## PROJECTS

---

### XOdia 2019: AI vs AI Competition Platform | [github.com/rhn19/Xodia\\_2019](https://github.com/rhn19/Xodia_2019)

- Enabled execution of C++ bot matches within resource-limited Docker environments using Django with threaded producer-consumer queues and reCAPTCHA-protected user models.
- Optimized search heuristics for sub-second decision-making in Abalone ( $10^{154}$  complexity) by engineering C++ agents utilizing MiniMax with Alpha-Beta pruning.

### IndicLM: Language Models for Indian Languages | [github.com/rhn19/indic-LM](https://github.com/rhn19/indic-LM)

- Developed tokenizers for 10+ Indic languages by training BPE, WordPiece, and SentencePiece on 10M+ IndicNLP and Wikipedia samples using PyTorch and Hugging Face Transformers.
- Attained 37.06 perplexity and 96% accuracy on downstream tasks for Marathi by training an ALBERT model via Masked Language Modeling.

### TCGuard: Legal Clause Analysis using Deep Learning | [github.com/rhn19/devhack](https://github.com/rhn19/devhack)

- Identified risky EULA clauses with 96% accuracy by training a PyTorch Bi-LSTM on 23,000 legal documents and deploying the model via Flask.

### High-Performance Entity Simulation (Unity DOTS/ECS) | [github.com/rhn19/ECS-Zombies](https://github.com/rhn19/ECS-Zombies)

- Simulated 5,000+ autonomous agents at high frame rates by engineering a Unity ECS system utilizing Burst Compiler, parallel job scheduling, and cache-friendly memory.

## TECHNICAL SKILLS

---

**Languages:** C++, Python, C#, SQL, Bash, GLSL

**AI/ML:** PyTorch, ExecuTorch, TorchAO, HuggingFace Transformers, LangChain

**XR/Graphics:** OpenXR, Unity, Unreal Engine, DirectX, OpenGL, Vulkan, ARFoundation

**Tools:** Git, Docker, Azure DevOps, OpenCV, Django, Flask, WiX Toolset, Maya, Visual Studio, Android Studio, Xcode

**Platforms:** iOS, Android, Windows, Linux, Web, AWS, Spatial Reality Display

**Expertise:** On-Device/Edge AI, LLM Fine-tuning/Quantization/Pruning/Distillation, GenAI (RAG), AR/VR/XR, 3D Math, Multi-threading, NLP